

Figure 4 — Cross-system generalisation (train on all, eval per type)

